

# Practical Deep learning for Computer Vision

---

Dr Kien Nguyen

Email: [k.nguyenthanh@qut.edu.au](mailto:k.nguyenthanh@qut.edu.au)

Queensland University of Technology, Australia

## **Warning**

All materials belong to Dr Kien Nguyen and copyright-protected and law-protected. Any sharing has to be inferred to Dr Kien Nguyen.

# Deep learning is new SEXY (for a reason!)



## **Week 3: Object Recognition**

Resnet from scratch

Transfer Learning

Google Colab

## **Week 4: Object Detection**

DarkNet & YOLO

Transfer Learning

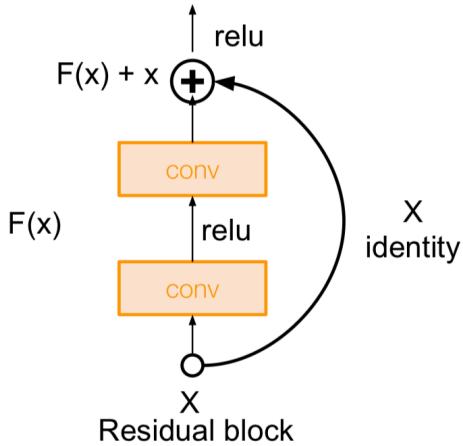
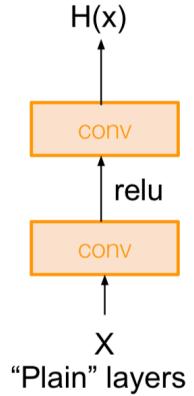
AWS

## **Week 5: GAN**

SSD

GAN

# Week 3. Resnet from scratch





[www.image-net.org](http://www.image-net.org)

**22K** categories and **14M** images

- Animals
  - Bird
  - Fish
  - Mammal
  - Invertebrate
- Plants
  - Tree
  - Flower
  - Food
  - Materials
- Structures
  - Artifact
  - Tools
  - Appliances
  - Structures
- Person
- Scenes
  - Indoor
  - Geological Formations
- Sport Activities



Deng, Dong, Socher, Li, Li, & Fei-Fei, 2009

# IMAGENET Large Scale Visual Recognition Challenge

Steel drum

The Image Classification Challenge:  
1,000 object classes  
1,431,167 images



Output:  
Scale  
T-shirt  
Steel drum  
Drumstick  
Mud turtle

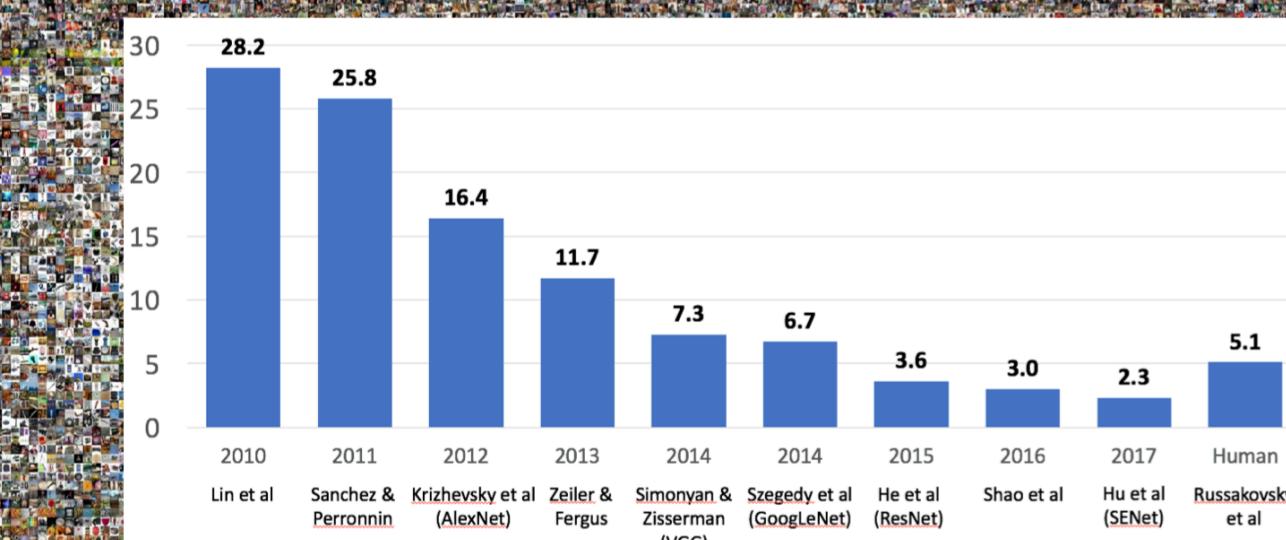


Output:  
Scale  
T-shirt  
Giant panda  
Drumstick  
Mud turtle



# IMAGENET Large Scale Visual Recognition Challenge

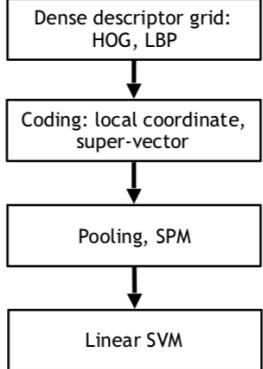
The Image Classification Challenge:  
1,000 object classes  
1,431,167 images



# IMAGENET Large Scale Visual Recognition Challenge

## Year 2010

NEC-UIUC

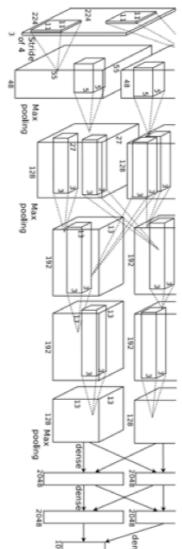


[Lin CVPR 2011]

Lion image by Swissfrog is licensed under CC BY 3.0

## Year 2012

SuperVision



[Krizhevsky NIPS 2012]

Figure copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012. Reproduced with permission.

## Year 2014

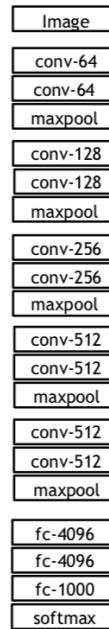
GoogLeNet

- Pooling
- Convolution
- Softmax
- Other



[Szegedy arxiv 2014]

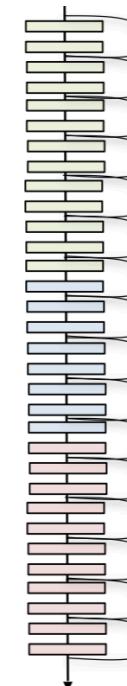
VGG



[Simonyan arxiv 2014]

## Year 2015

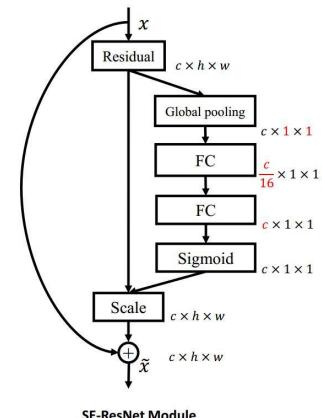
ResNet



[He ICCV 2015]

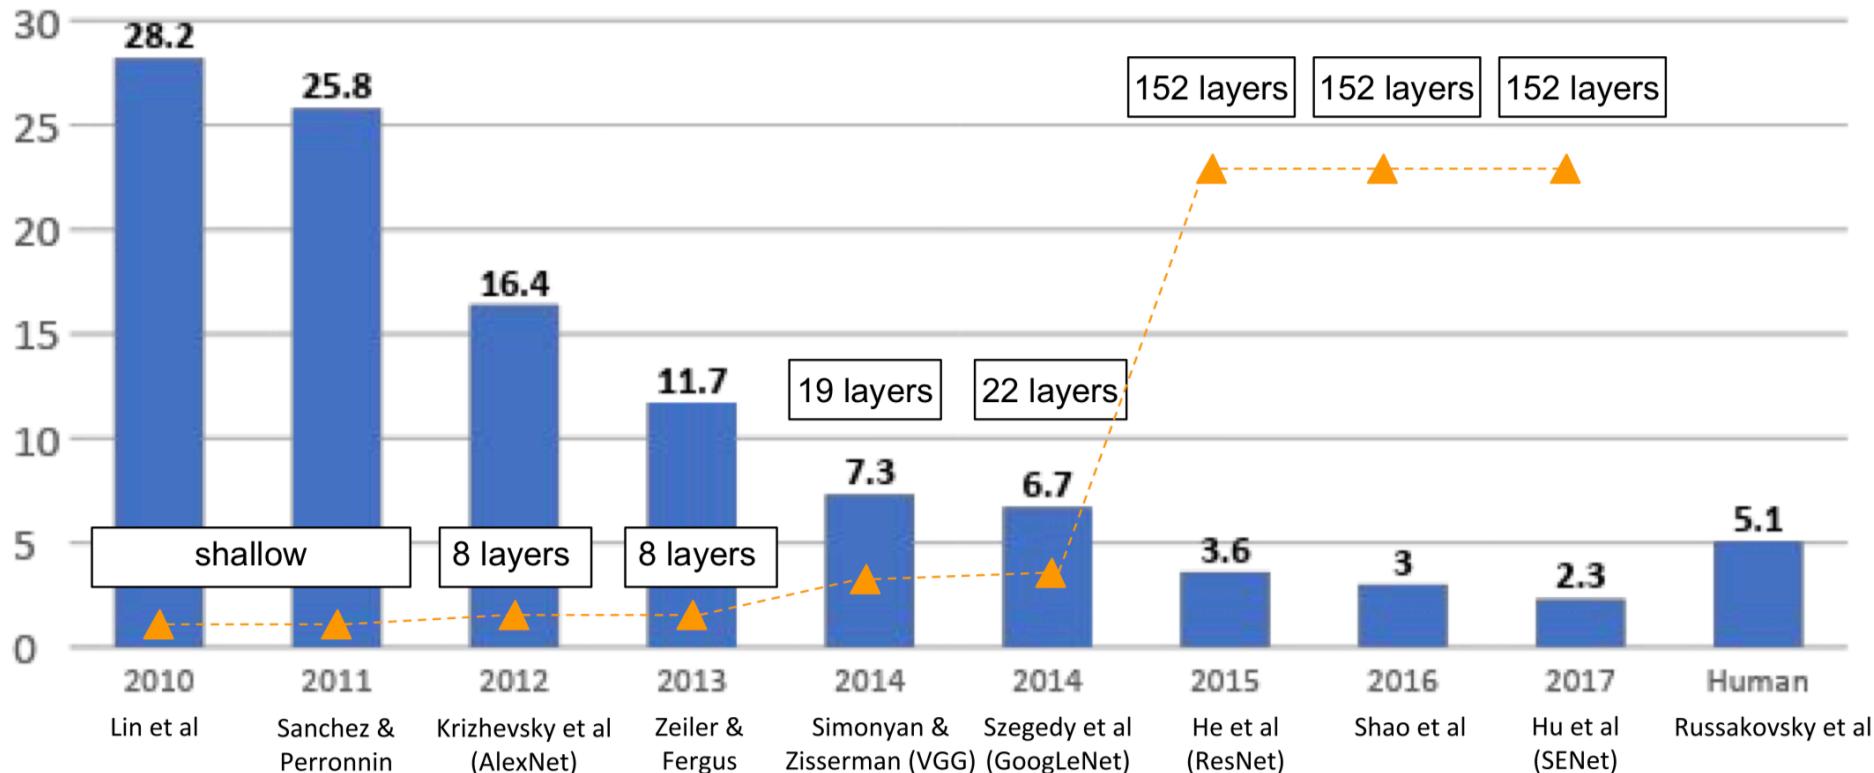
## Year 2017

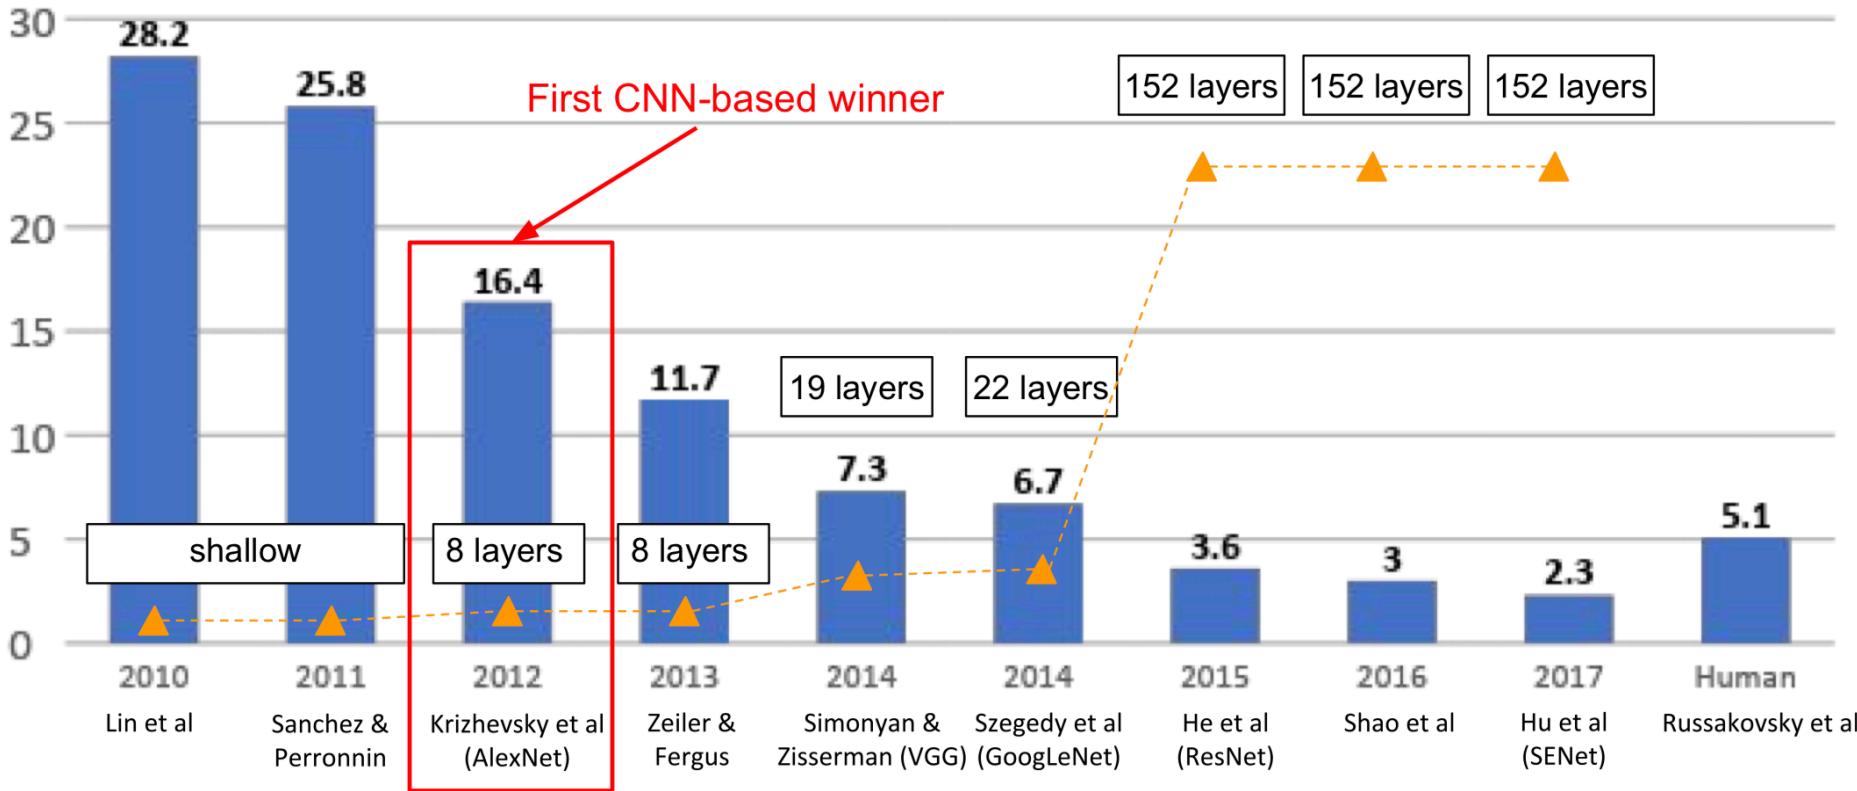
SENet

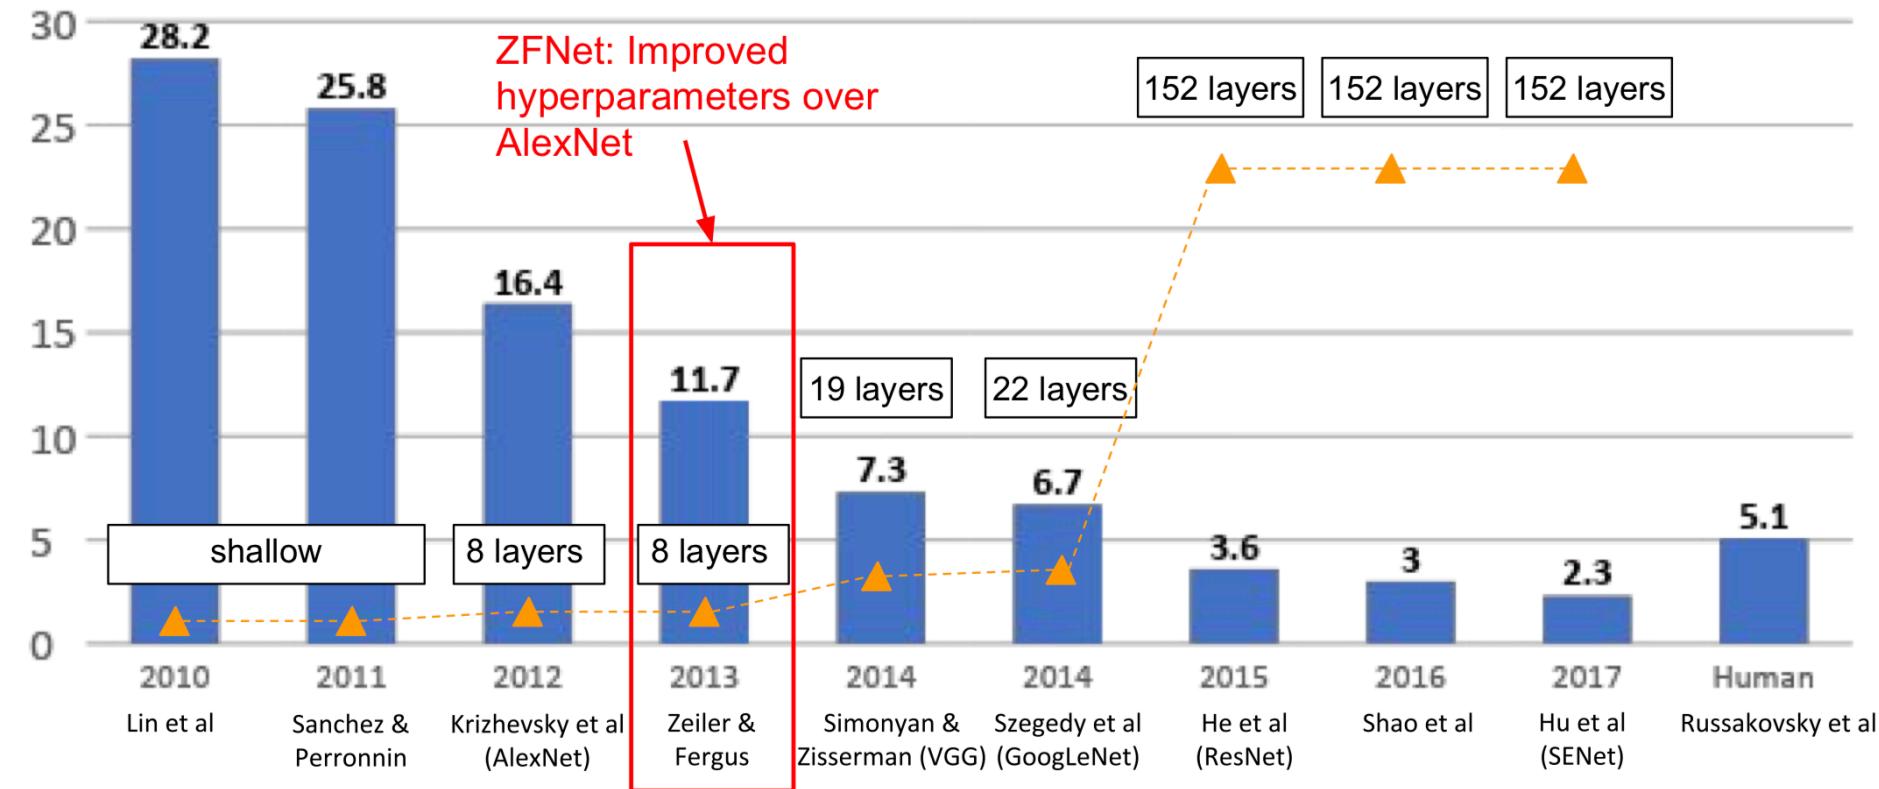


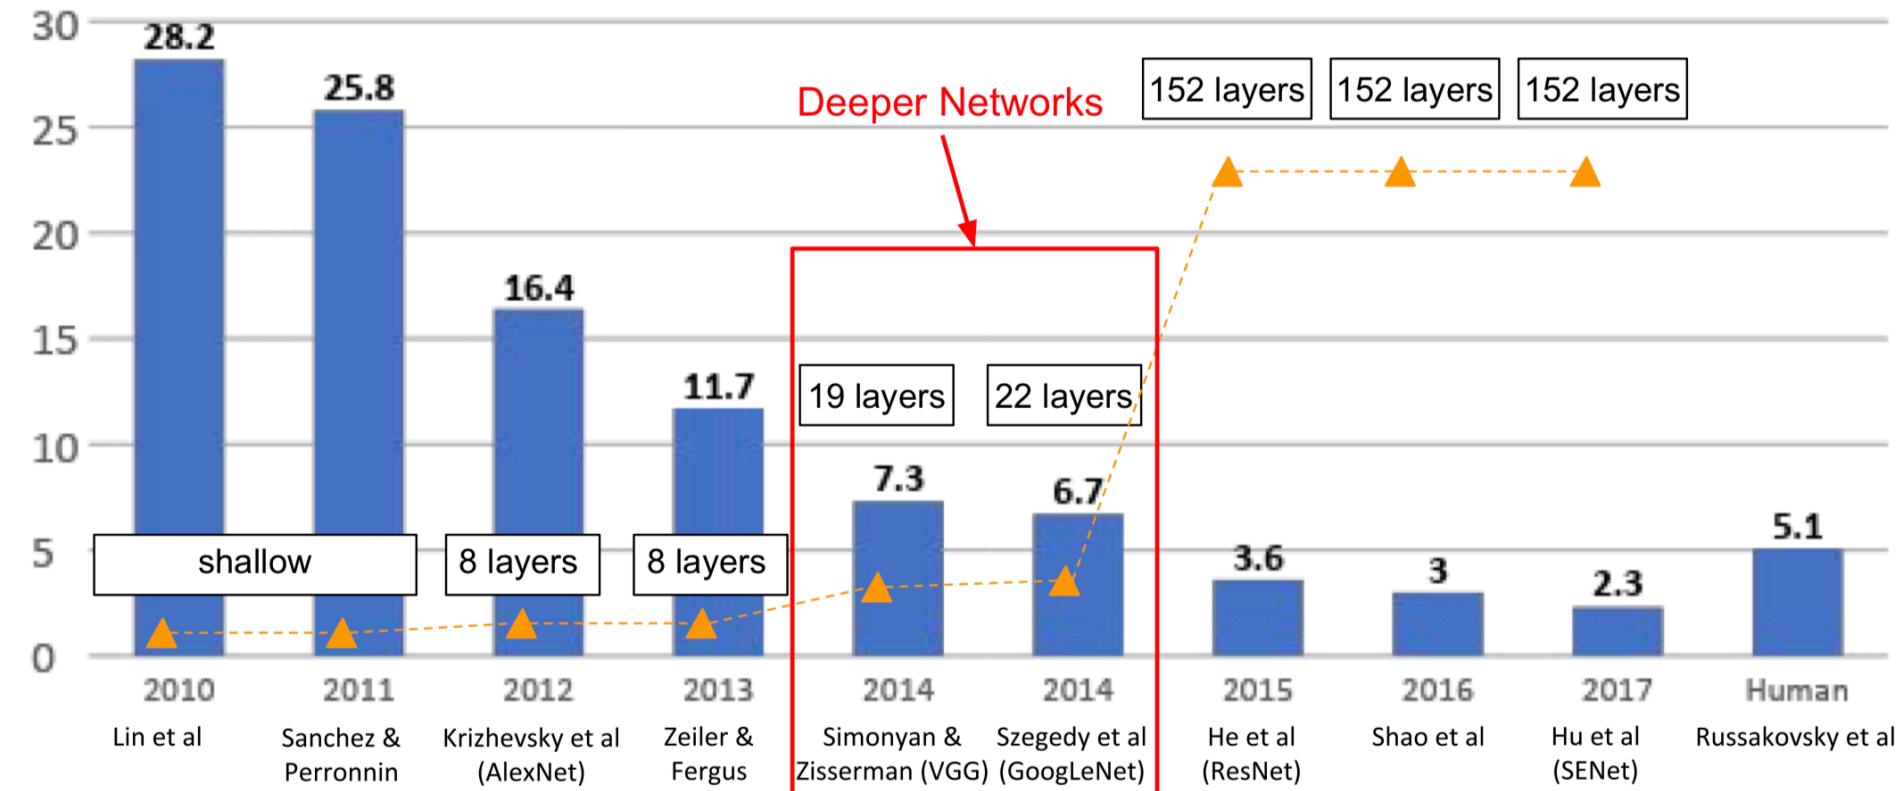
[Hu CVPR 2018]

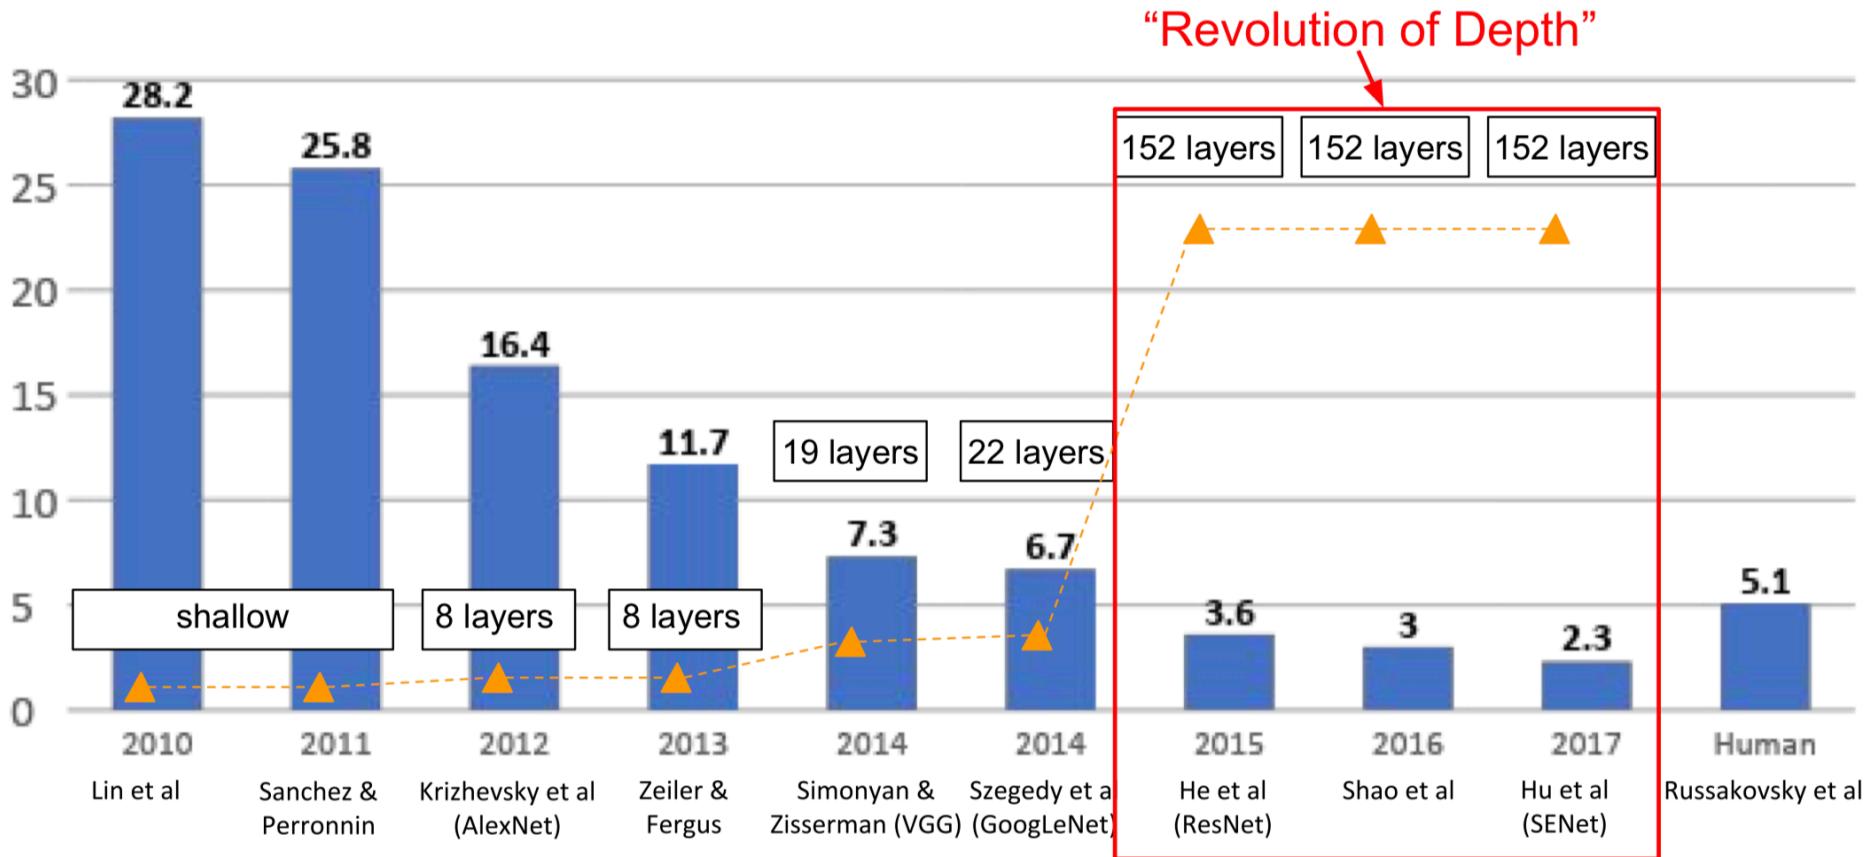
# Revolution of Depth



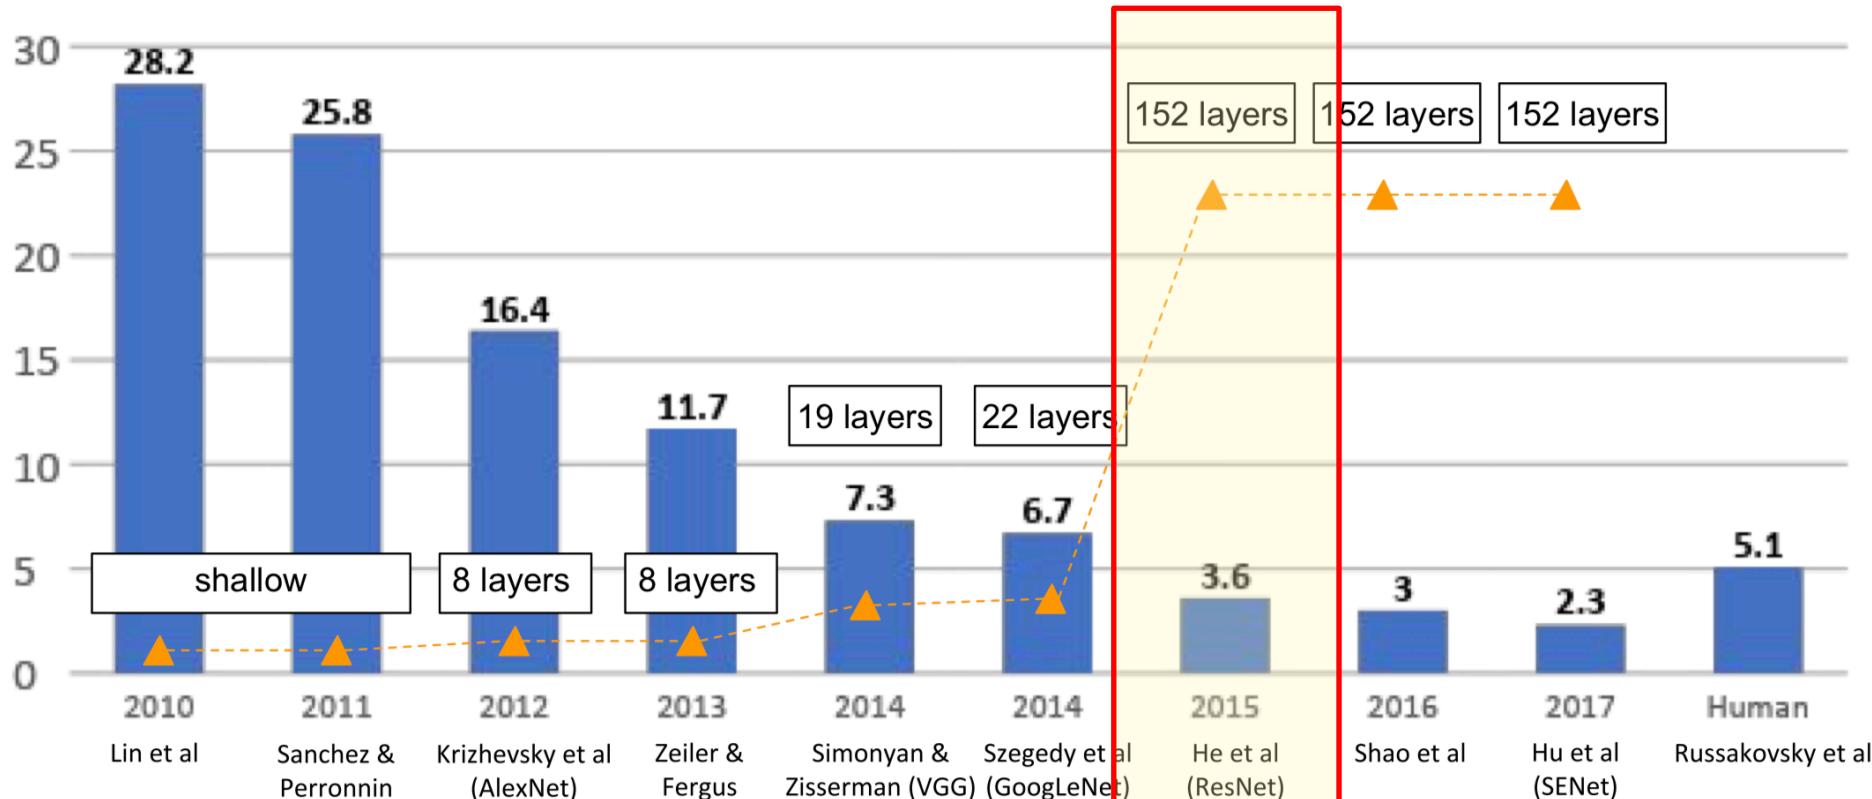








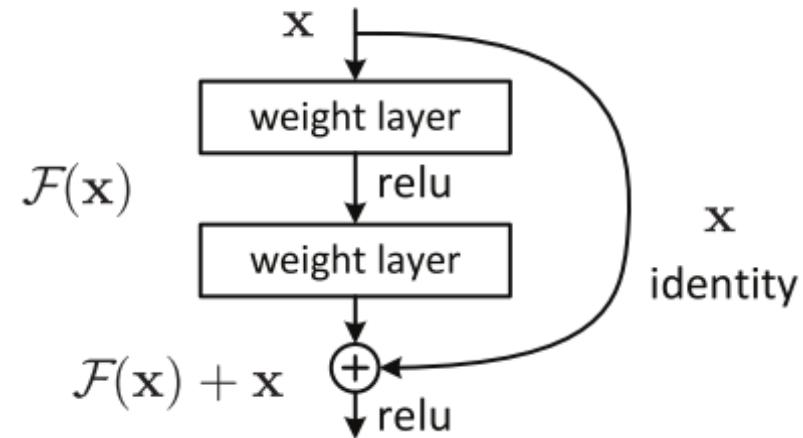
# Revolution of Depth



# Microsoft's ResNet

[He et al., 2015]

- WHAT?
  - Very deep networks using residual connections
- 152-layer model for ImageNet
- ILSVRC'15 classification winner  
(3.57% top 5 error)
- Swept all classification and detection competitions in ILSVRC'15 and COCO'15!

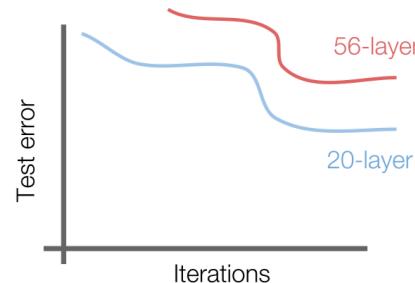
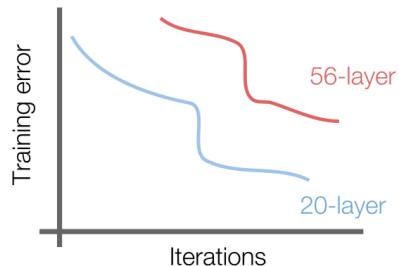


# Microsoft's ResNet

[He et al., 2015]

- WHY?

- What happens when we continue stacking deeper layers on a “plain” CNN?



56-layer model performs worse on both training and test error  
-> The deeper model performs worse, but it's not caused by overfitting!

Hypothesis:

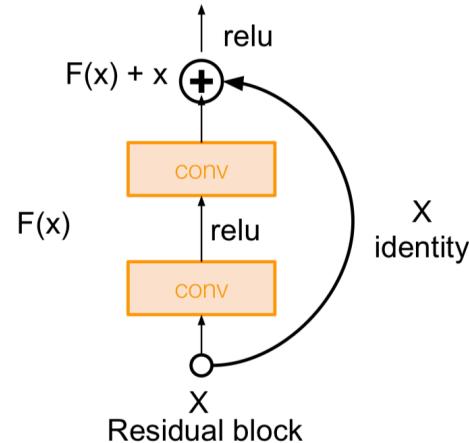
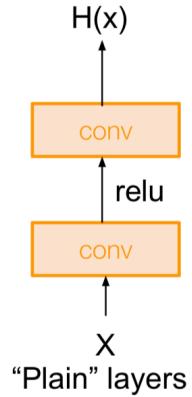
- the problem is an *optimization* problem, deeper models are harder to optimize
- Harder because the gradient vanishing and exploding effect

# Microsoft's ResNet

[He et al., 2015]

- **HOW?**

- Use network layers to fit a residual mapping instead of directly trying to fit a desired underlying mapping

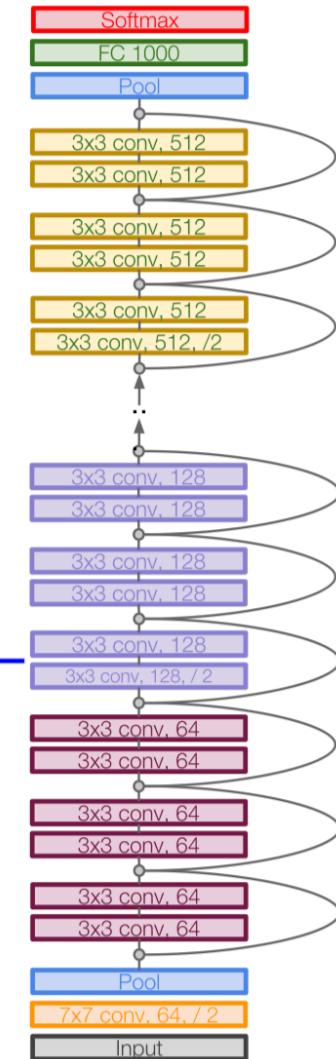
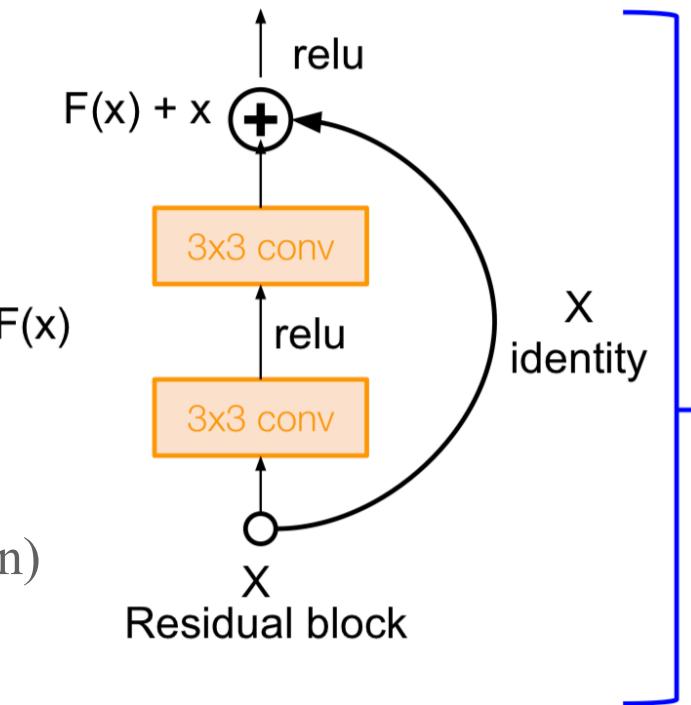


# Microsoft’s ResNet

[He et al., 2015]

## Full ResNet architecture:

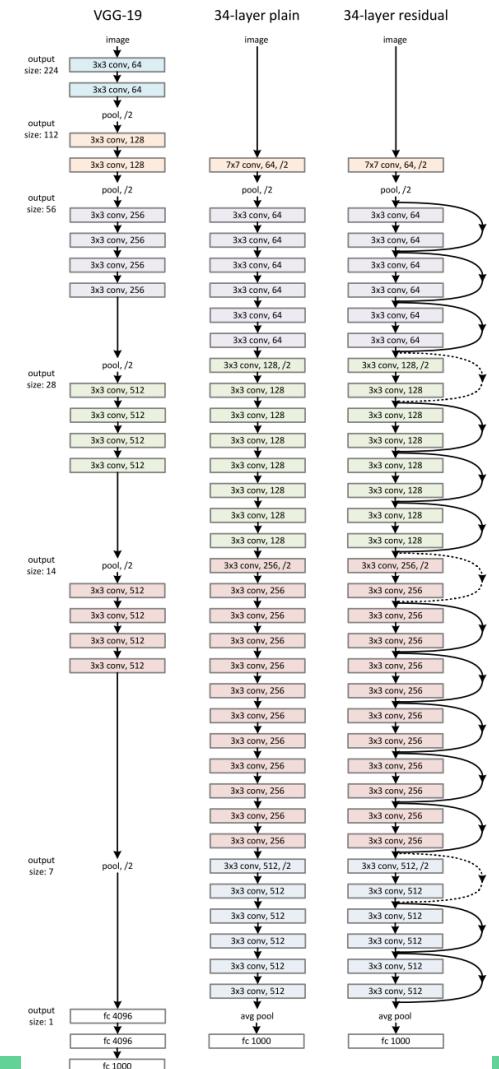
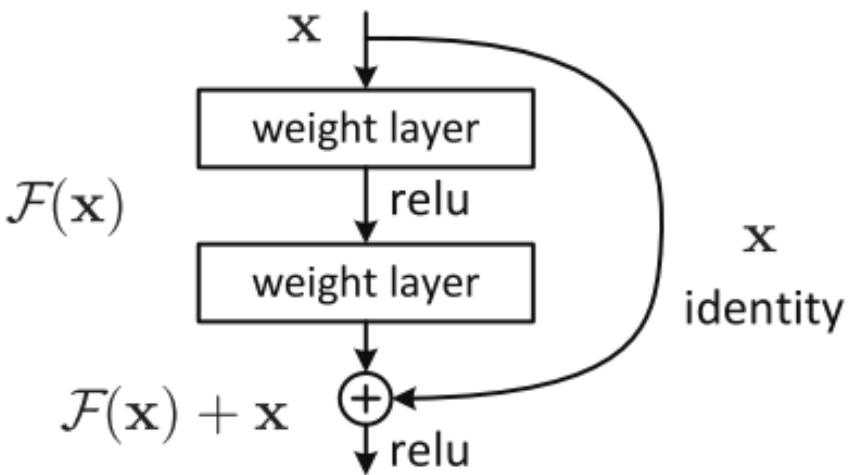
- Stack residual blocks
  - Every residual block has two 3x3 conv layers
  - Periodically, double # of filters and downsample spatially using stride 2 (/2 in each dimension)



# Microsoft's ResNet

[He et al., 2015]

- Very deep networks using residual connections



# ResNet variants

layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112			7×7, 64, stride 2		
conv2_x	56×56	$\begin{bmatrix} 3\times3, 64 \\ 3\times3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3\times3, 64 \\ 3\times3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times1, 64 \\ 3\times3, 64 \\ 1\times1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times1, 64 \\ 3\times3, 64 \\ 1\times1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times1, 64 \\ 3\times3, 64 \\ 1\times1, 256 \end{bmatrix} \times 3$
conv3_x	28×28	$\begin{bmatrix} 3\times3, 128 \\ 3\times3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3\times3, 128 \\ 3\times3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1\times1, 128 \\ 3\times3, 128 \\ 1\times1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1\times1, 128 \\ 3\times3, 128 \\ 1\times1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1\times1, 128 \\ 3\times3, 128 \\ 1\times1, 512 \end{bmatrix} \times 8$
conv4_x	14×14	$\begin{bmatrix} 3\times3, 256 \\ 3\times3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3\times3, 256 \\ 3\times3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1\times1, 256 \\ 3\times3, 256 \\ 1\times1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1\times1, 256 \\ 3\times3, 256 \\ 1\times1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1\times1, 256 \\ 3\times3, 256 \\ 1\times1, 1024 \end{bmatrix} \times 36$
conv5_x	7×7	$\begin{bmatrix} 3\times3, 512 \\ 3\times3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3\times3, 512 \\ 3\times3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times1, 512 \\ 3\times3, 512 \\ 1\times1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times1, 512 \\ 3\times3, 512 \\ 1\times1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times1, 512 \\ 3\times3, 512 \\ 1\times1, 2048 \end{bmatrix} \times 3$
	1×1	average pool, 1000-d fc, softmax				
FLOPs		$1.8 \times 10^9$	$3.6 \times 10^9$	$3.8 \times 10^9$	$7.6 \times 10^9$	$11.3 \times 10^9$

# Training ResNet in practice

- Batch Normalization after every CONV layer
- Xavier 2/ initialization from He et al.
- SGD + Momentum (0.9)
- Learning rate: 0.1, divided by 10 when validation error plateaus
- Mini-batch size 256
- Weight decay of 1e-5
- No dropout used

# Implementations on popular platforms

- Tensorflow implementation
  - <https://github.com/tensorflow/tensorflow/tree/361a82d73a50a800510674b3aaa20e4845e56434/tensorflow/contrib/slim/python/slim/nets>
- Pytorch implementation
  - <https://github.com/pytorch/vision/blob/master/torchvision/models/resnet.py>

# Experimental Results

- Able to train very deep networks without degrading (152 layers on ImageNet, 1202 on Cifar)
- Deeper networks now achieve lowing training error as expected
- Swept 1st place in all ILSVRC and COCO 2015 competitions

## MSRA @ ILSVRC & COCO 2015 Competitions

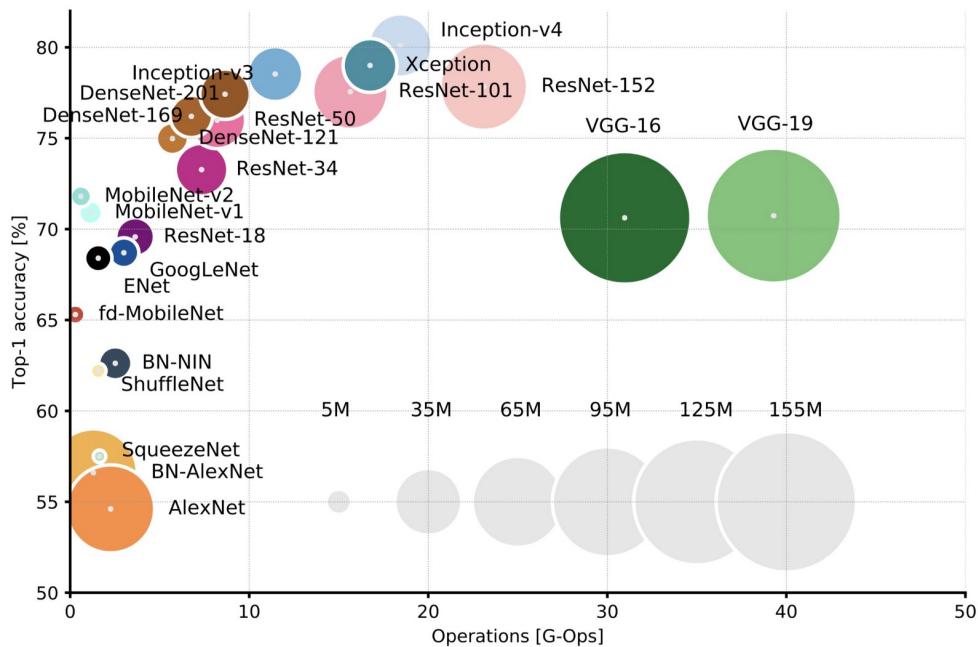
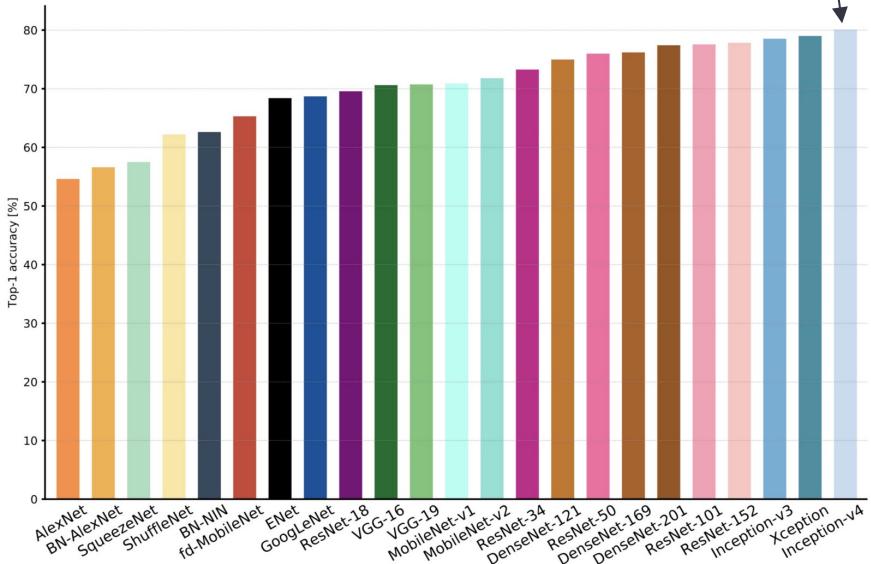
- **1st places in all five main tracks**

- ImageNet Classification: “Ultra-deep” (quote Yann) **152-layer** nets
- ImageNet Detection: **16%** better than 2nd
- ImageNet Localization: **27%** better than 2nd
- COCO Detection: **11%** better than 2nd
- COCO Segmentation: **12%** better than 2nd

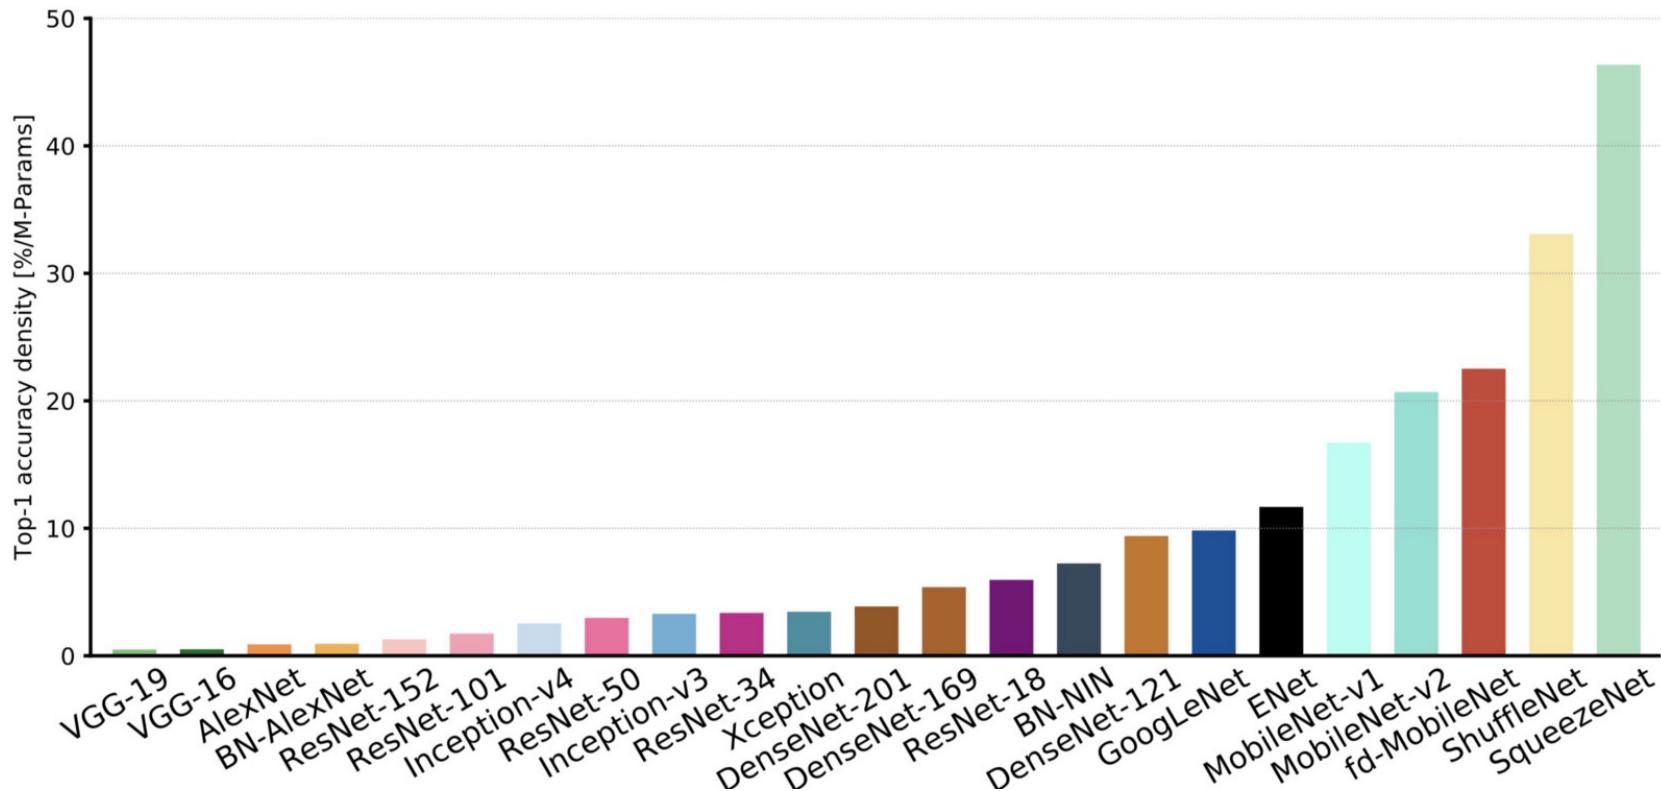
ILSVRC 2015 classification winner  
(3.6% top 5 error) -- better than “human performance”! (Russakovsky 2014)

# Top1 vs. operations vs. parameters

Inception-v4: Resnet + Inception!



# Accuracy per parameter



# Transfer Learning

- **WHAT?**

- Instead, it is common to pretrain a ConvNet on a very large dataset (e.g. ImageNet, which contains 1.2 million images with 1000 categories), then apply Transfer Learning to adapt the pretrained model to a new task

- **WHY?**

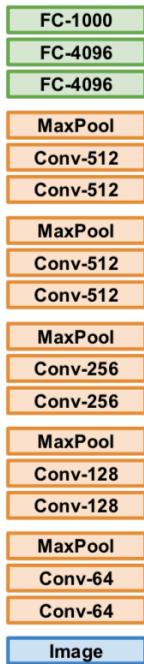
- Computation: modern ConvNets take 2-3 weeks to train across multiple GPUs
- Data: on ImageNet (>1M images)
- => Rarely you have these luxury resources available

# Transfer Learning

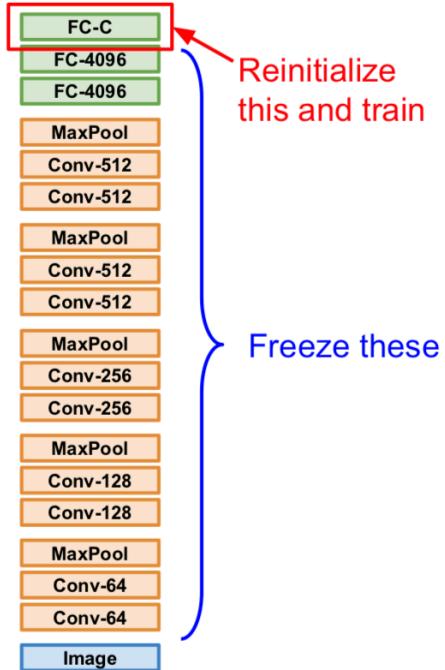
- **HOW?**
  - ConvNet as fixed feature extractor
  - Fine-tuning the ConvNet

# Transfer Learning

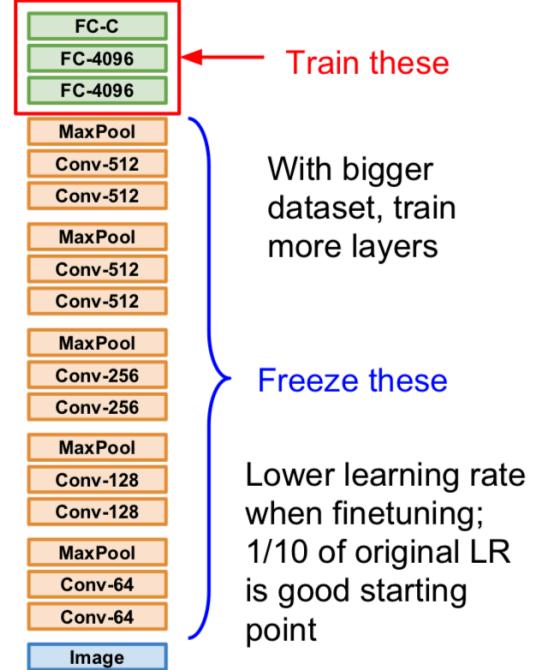
1. Train on Imagenet



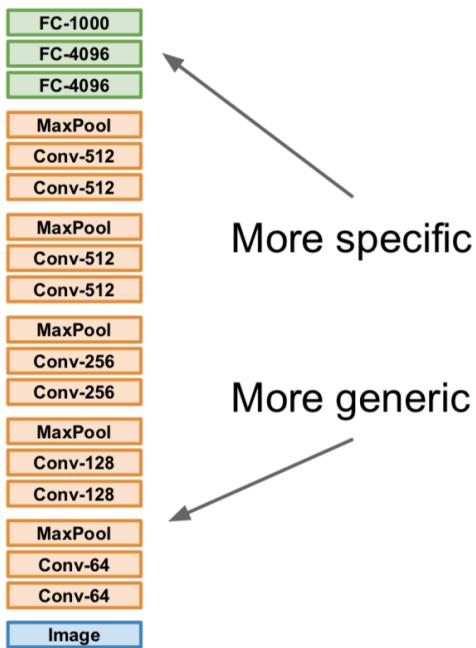
2. Small Dataset (C classes)



3. Bigger dataset



# Transfer Learning



	<b>very similar dataset</b>	<b>very different dataset</b>
<b>very little data</b>	Use Linear Classifier on top layer	You're in trouble... Try linear classifier from different stages
<b>quite a lot of data</b>	Finetune a few layers	Finetune a larger number of layers

# Google Colab

- Google Colab is a *free cloud* service and now it supports free *GPU*!
- You can;
  - improve your **Python** programming language coding skills.
  - develop deep learning applications using popular libraries such as **Keras**, **TensorFlow**, **PyTorch**, **OpenCV**

# Google Colab

- Tutorial
  - <https://medium.com/deep-learning-turkey/google-colab-free-gpu-tutorial-e113627b9f5d>
  - **Creating Folder on Google Drive**
  - **Creating New Colab Notebook**
  - **Setting Python version and Free CPU/GPU**
  - **Running Basic Python Codes with Google Colab**
  - Accessing data from Google Drive

# Resnet and Transfer learning tutorial codes on Colab

- Design your own simple CNN
- Use pretrained Resnets for classification
- Visualize Resnet networks
- Design your own ResNet from scratch
- Transfer learning with Resnet
- All my tutorial codes are here:
  - Keras:
    - <https://colab.research.google.com/drive/1iVChsQa7dAJych6gciq21QZB5tEVlzIw>
    - <https://colab.research.google.com/drive/159QWa2LUIf8cQQ2ksnKVjBYnEb6wNKb3>
  - Pytorch: <https://colab.research.google.com/drive/1MoE5CvDz5pOOoJkmWRstqqb1JzD46Lmt>